

# 【全人類がわかるE資格コース PyTorch版】 プロダクト開発演習

## 6. Labolo Tomato事前学習モデルの検証

公開されているLabolo Tomato Datasetにより学習済みのMask R-CNNモデルが実際に動作することを確認する。google Colab環境でGPUを使用する。モデル環境の構築の後、test data、アナリーションデータを使って公開されている性能がでるのかを評価し、その正当性を確認(validation)する。

更に、今回新たに準備したトマト画像（静止画、動画）データによって、正しく判別が行えるかを検証（verification）する。

```
In [1]: # Google Driveのマウント
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

## 6.1 mmdetection v2.23.0 のインストール (google Colab版)

[https://github.com/open-mmlab/mmdetection/blob/master/demo/MMDet\\_Tutorial.ipynb](https://github.com/open-mmlab/mmdetection/blob/master/demo/MMDet_Tutorial.ipynb)

```
In [2]: # pritrainモデルの動作確認
# MMDetectionのインストール
# Check nvcc version
!nvcc -V
# Check GCC version
!gcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2020 NVIDIA Corporation
Built on Mon_Oct_12_20:09:46_PDT_2020
Cuda compilation tools, release 11.1, V11.1.105
Build cuda_11.1.1.TC455_06.29190527_0
gcc (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0
Copyright (C) 2017 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
In [3]: # ***** こちらは旧バージョンをロード v2.24.0～最新版v2.25.0 では学習(tools/t
# 関係ファイルのインストール(旧バージョン用)
# install dependencies: (use cu111 because colab has CUDA 11.1)
!pip install torch==1.9.0+cu111 torchvision==0.10.0+cu111 -f https://download.pytorch.org
# install mmcv-full thus we could use CUDA operators
### !pip install mmcv-full -f https://download.openmmlab.com/mmcv/dist/cu111/torch1.9.
### バージョンv2.25.0ではAttributeError: 'ConfigDict' object has no attribute 'device'
### mmdetection v.2.23.0を使用するため、それに合わせてmmcvも1.5.x から1.3.17にダウンロード
!pip install mmcv-full==1.3.17 -f https://download.openmmlab.com/mmcv/dist/cu111/torch1.9.0+cu111
# Install mmdetection
!rm -rf mmdetection
```

```
### !git clone https://github.com/open-mmlab/mmdetection.git
### mmdetection 最新バージョンv2.25.0ではtools/train.pyに問題があり、
### AttributeError: 'ConfigDict' object has no attribute 'device' が発生するので、回避策
!git clone https://github.com/open-mmlab/mmdetection.git -b v2.23.0 --depth 1    ### 1
%cd mmdetection

!pip install -e .
```

```

that are installed. This behaviour is the source of the following dependency conflict
s.
torchtext 0.12.0 requires torch==1.11.0, but you have torch 1.9.0+cu111 which is incom
patible.
torchaudio 0.11.0+cu113 requires torch==1.11.0, but you have torch 1.9.0+cu111 which i
s incompatible.
Successfully installed torch-1.9.0+cu111 torchvision-0.10.0+cu111
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pu
blic/simple/
Looking in links: https://download.openmmlab.com/mmcv/dist/cu111/torch1.9.0/index.html
Collecting mmcv-full==1.3.17
  Downloading https://download.openmmlab.com/mmcv/dist/cu111/torch1.9.0/mmcv_full-1.3.
17-cp37-cp37m-manylinux1_x86_64.whl (50.4 MB)
    |████████████████████████████████| 50.4 MB 9.3 MB/s
Requirement already satisfied: opencv-python>=3 in /usr/local/lib/python3.7/dist-pac
kages (from mmcv-full==1.3.17) (4.1.2.30)
Requirement already satisfied: Pillow in /usr/local/lib/python3.7/dist-packages (from
mmcv-full==1.3.17) (7.1.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.7/dist-packages (fr
om mmcv-full==1.3.17) (21.3)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from m
mcv-full==1.3.17) (1.21.6)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.7/dist-packages (from
mmcv-full==1.3.17) (3.13)
Collecting yapf
  Downloading yapf-0.32.0-py2.py3-none-any.whl (190 kB)
    |████████████████████████████████| 190 kB 5.1 MB/s
Collecting addict
  Downloading addict-2.4.0-py3-none-any.whl (3.8 kB)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.7/di
st-packages (from packaging->mmcv-full==1.3.17) (3.0.9)
Installing collected packages: yapf, addict, mmcv-full
Successfully installed addict-2.4.0 mmcv-full-1.3.17 yapf-0.32.0
Cloning into 'mmdetection'...
remote: Enumerating objects: 1575, done.
remote: Counting objects: 100% (1575/1575), done.
remote: Compressing objects: 100% (1094/1094), done.
remote: Total 1575 (delta 612), reused 781 (delta 466), pack-reused 0
Receiving objects: 100% (1575/1575), 16.33 MiB | 30.57 MiB/s, done.
Resolving deltas: 100% (612/612), done.
Note: checking out '3e2693151add9b5d6db99b944da020cba837266b'.
```

You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by performing another checkout.

If you want to create a new branch to retain commits you create, you may do so (now or later) by using -b with the checkout command again. Example:

```

git checkout -b <new-branch-name>

/content/mmdetection
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pu
blic/simple/
Obtaining file:///content/mmdetection
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (f
rom mmdet==2.23.0) (3.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from m
mdet==2.23.0) (1.21.6)
Requirement already satisfied: pycocotools in /usr/local/lib/python3.7/dist-packages
(from mmdet==2.23.0) (2.0.4)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from mmd
et==2.23.0) (1.15.0)
Collecting terminaltables
```

```
Downloading terminaltables-3.1.10-py3.py3-none-any.whl (15 kB)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->mmdet==2.23.0) (3.0.9)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->mmdet==2.23.0) (1.4.3)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->mmdet==2.23.0) (2.8.2)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib->mmdet==2.23.0) (0.11.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from kiwisolver>=1.0.1->matplotlib->mmdet==2.23.0) (4.1.1)
Installing collected packages: terminaltables, mmdet
  Running setup.py develop for mmdet
Successfully installed mmdet-2.23.0 terminaltables-3.1.10
```

In [4]:

```
# インストール環境の確認
from mmcv import collect_env
collect_env()
```

Out[4]:

```
{'CUDA available': True,
 'CUDA_HOME': '/usr/local/cuda',
 'GCC': 'gcc (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0',
 'GPU 0': 'Tesla P100-PCIE-16GB',
 'MMCV': '1.3.17',
 'MMCV CUDA Compiler': '11.1',
 'MMCV Compiler': 'GCC 7.3',
 'NVCC': 'Build cuda_11.1.TC455_06.29190527_0',
 'OpenCV': '4.1.2',
 'PyTorch': '1.9.0+cu111',
 'PyTorch compiling details': 'PyTorch built with:¥n - GCC 7.3¥n - C++ Version: 201402¥n - Intel(R) Math Kernel Library Version 2020.0.0 Product Build 20191122 for Intel(R) 64 architecture applications¥n - Intel(R) MKL-DNN v2.1.2 (Git Hash 98be7e8afa711dc9b66c8ff3504129cb82013cdb)¥n - OpenMP 201511 (a.k.a. OpenMP 4.5)¥n - NNPACK is enabled¥n - CPU capability usage: AVX2¥n - CUDA Runtime 11.1¥n - NVCC architecture flags: -gencode;arch=compute_37,code=sm_37;-gencode;arch=compute_50,code=sm_50;-gencode;arch=compute_60,code=sm_60;-gencode;arch=compute_70,code=sm_70;-gencode;arch=compute_75,code=sm_75;-gencode;arch=compute_80,code=sm_80;-gencode;arch=compute_86,code=sm_86¥n - CuDNN 8.0.5¥n - Magma 2.5.2¥n - Build settings: BLAS_INFO=mkl, BUILD_TYPE=Release, CUDA_VERSION=11.1, CUDNN_VERSION=8.0.5, CXX_COMPILER=/opt/rh/devtoolset-7/root/usr/bin/c++, CXX_FLAGS= -Wno-deprecated -fvisibility-inlines-hidden -DUSE_PTHREADPOOL -fopenmp -DNDEBUG -DUSE_KINET0 -DUSE_FBGEMM -DUSE_QNNPACK -DUSE_PYTORCH_QNNPACK -DUSE_XNNPACK -DSYMBOLICATE_MOBILE_DEBUG_HANDLE -O2 -fPIC -Wno-narrowing -Wall -Wextra -Werror=return-type -Wno-missing-field-initializers -Wno-type-limits -Wno-array-bounds -Wno-unknown-pragmas -Wno-sign-compare -Wno-unused-parameter -Wno-unused-variable -Wno-unused-function -Wno-unused-result -Wno-unused-local-typedefs -Wno-strict-overflow -Wno-strict-aliasing -Wno-error=deprecated-declarations -Wno-stringop-overflow -Wno-psabi -Wno-error=pedantic -Wno-error=redundant-decls -Wno-error=old-style-cast -fdiagnostics-color=always -faligned-new -Wno-unused-but-set-variable -Wno-maybe-uninitialized -fno-math-errno -fno-trapping-math -Werror=format -Wno-stringop-overflow, LAPACK_INFO=mkl, PERF_WITH_AVX=1, PERF_WITH_AVX2=1, PERF_WITH_AVX512=1, TORCH_VERSION=1.9.0, USE_CUDA=ON, USE_CUDNN=ON, USE_EXCEPTION_PTR=1, USE_GFLAGS=OFF, USE_GLOG=OFF, USE_MKL=ON, USE_MKLDNN=0N, USE_MPI=OFF, USE_NCCL=ON, USE_NNPACK=ON, USE_OPENMP=ON, ¥n',
 'Python': '3.7.13 (default, Apr 24 2022, 01:04:09) [GCC 7.5.0]',
 'TorchVision': '0.10.0+cu111',
 'sys.platform': 'linux'}
```

In [5]:

```
# Check Pytorch installation
import torch, torchvision
print(torch.__version__, torch.cuda.is_available())

# Check MMDetection installation
import mmdet
print(mmdet.__version__)
```

```
# Check mmcv installation
from mmcv.ops import get_compiling_cuda_version, get_compiler_version
print(get_compiling_cuda_version())
print(get_compiler_version())
```

```
1. 9. 0+cu111 True
2. 23. 0
11. 1
GCC 7. 3
```

## 6.2 LABORO TOMATOデータ、学習済みモデルの移植、動作確認

### Mask R-CNN r50 FPN 1x8 COCOモデル

Laboro Tomato dataset (ローカルコピーしたもの) を、google DriveからMMdetectionの実行環境にコピー（シンボリックリンク）

MMDetectionフレームワークで利用できるようモデルの環境設定(config) を、LaboroTomatoの公開情報に基づき変更する。

- Datasetパス情報の変更
- 識別クラス数の変更 (80 => 6)
- 学習済みcheckpointファイル（パラメタデータファイル）のロードである。

<https://github.com/open-mmlab/mmdetection>

```
In [6]: !cd '/content/mmdetection'

# laboro_tomato datasetを./data/にシンボリックリンク
!mkdir -p ./data
!ln -s '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato' './data/'
```

```
In [7]: # datasetクラスの登録
!ln -s '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato/etc/laboro_tomato.py' './mmdet/datasets/laboro_tomato.py'
```

```
In [8]: # configのセット
from mmcv import Config
cfg = Config.fromfile('./configs/mask_rcnn/mask_rcnn_r50_fpn_1x_coco.py')

# 学習済みcheckpointを./checkpoints/にシンボリックリンク
!mkdir -p ./checkpoints
!ln -s '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato/etc/laboro_tomato_48ep.pth' './checkpoints/laboro_tomato_48ep.pth'
```

```
In [9]: # Add dataset names to mmdet/datasets/__init__.py
!cp '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato/etc/__init__.py' './mmdet/datasets/__init__.py'

# Configuration files setup
# step1: add laboro_tomato_base.py to configs/_base_/datasets/
!cp '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato/etc/laboro_tomato_base.py' './configs/_base_/datasets/laboro_tomato_base.py'
```

```
# step2: add laboro_tomato_instance.py to configs/_base_/datasets/
!cp '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato/etc/laboro_tomato_instance.py' './configs/_base_/datasets/laboro_tomato_instance.py'

# step3: overwrite class numbers at model configuration file configs/_base_/models/mask_rcnn_r50_fpn.py
!cp '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato/etc/mask_rcnn_r50_fpn.py' './configs/_base_/models/mask_rcnn_r50_fpn.py'

# step4: overwrite base instance configuration file name at configs/mask_rcnn/mask_rcnn_r50_fpn_1x_coco.py
!cp '/content/drive/MyDrive/Colab Notebooks/product_develop/laboro_tomato/etc/mask_rcnn_r50_fpn_1x_coco.py' './configs/mask_rcnn/mask_rcnn_r50_fpn_1x_coco.py'
```

In [10]:

```
# ネットワークモデルのコンフィグレーション
import mmcv
from mmcv.runner import load_checkpoint

from mmdet.apis import inference_detector, show_result_pyplot
from mmdet.models import build_detector

# Choose to use a config and initialize the detector
config = './configs/mask_rcnn/mask_rcnn_r50_fpn_1x_coco.py'
# Setup a checkpoint file to load
checkpoint = './checkpoints/laboro_tomato_48ep.pth'

# Set the device to be used for evaluation
### device='cuda:0'
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu') ##### for colab

# Load the config
config = mmcv.Config.fromfile(config)
# Set pretrained to be None since we do not need pretrained model here
config.model.pretrained = None

# Initialize the detector
model = build_detector(config.model)

# Load checkpoint
checkpoint = load_checkpoint(model, checkpoint, map_location=device)

# Set the classes of models for inference
model.CLASSES = checkpoint['meta']['CLASSES']

# We need to set the model's cfg for inference
model.cfg = config

# Convert the model to GPU
model.to(device)
# Convert the model into evaluation mode
model.eval()
```

load checkpoint from local path: ./checkpoints/laboro\_tomato\_48ep.pth

Out[10]:

```
MaskRCNN(
    backbone: ResNet(
        conv1: Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
        bn1: BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        relu: ReLU(inplace=True)
        maxpool: MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    )
    layer1: ResLayer(
```

```

(0): Bottleneck(
    (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(1): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (relu): ReLU(inplace=True)
)
(2): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (relu): ReLU(inplace=True)
)
)
(layer2): ResLayer(
    (0): Bottleneck(
        (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), b
ias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
            (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        )
    )
)

```

```

(1): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
                (relu): ReLU(inplace=True)
)
(2): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
                (relu): ReLU(inplace=True)
)
(3): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
                (relu): ReLU(inplace=True)
)
)
(layer3): ResLayer(
    (0): Bottleneck(
        (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), b
ias=False)
            (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
                (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
                    (relu): ReLU(inplace=True)
                    (downsample): Sequential(
                        (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
                        (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
                    )
)
    (1): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)

```

```

ias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
            (relu): ReLU(inplace=True)
        )
    (2): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
            (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
            (relu): ReLU(inplace=True)
        )
    (3): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
            (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
            (relu): ReLU(inplace=True)
        )
    (4): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
            (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
            (relu): ReLU(inplace=True)
        )
    (5): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
            (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
            (relu): ReLU(inplace=True)
        )
    )
(layer4): ResLayer(
    (0): Bottleneck(
        (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), b
ias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
            (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
        )
    )
    (1): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), b
ias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
        (relu): ReLU(inplace=True)
    )
)
init_cfg={'type': 'Pretrained', 'checkpoint': 'torchvision://resnet50'}
(neck): FPN(
    (lateral_convs): ModuleList(
        (0): ConvModule(
            (conv): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        )
        (1): ConvModule(
            (conv): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
        )
        (2): ConvModule(
            (conv): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
        )
        (3): ConvModule(
            (conv): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
        )
    )
    (fpn_convs): ModuleList(
        (0): ConvModule(
            (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        )
    )
)

```

```

(1): ConvModule(
    (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
(2): ConvModule(
    (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
(3): ConvModule(
    (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
)
)
init_cfg=[{'type': 'Xavier', 'layer': 'Conv2d', 'distribution': 'uniform'}]
(rpn_head): RPNHead(
    (loss_cls): CrossEntropyLoss(avg_non_ignore=False)
    (loss_bbox): L1Loss()
    (rpn_conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (rpn_cls): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
    (rpn_reg): Conv2d(256, 12, kernel_size=(1, 1), stride=(1, 1))
)
)
init_cfg=[{'type': 'Normal', 'layer': 'Conv2d', 'std': 0.01}]
(roi_head): StandardROIHead(
    (bbox_roi_extractor): SingleROIExtractor(
        (roi_layers): ModuleList(
            (0): RoIAlign(output_size=(7, 7), spatial_scale=0.25, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
            (1): RoIAlign(output_size=(7, 7), spatial_scale=0.125, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
            (2): RoIAlign(output_size=(7, 7), spatial_scale=0.0625, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
            (3): RoIAlign(output_size=(7, 7), spatial_scale=0.03125, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
        )
    )
    (bbox_head): Shared2FCBBoxHead(
        (loss_cls): CrossEntropyLoss(avg_non_ignore=False)
        (loss_bbox): L1Loss()
        (fc_cls): Linear(in_features=1024, out_features=7, bias=True)
        (fc_reg): Linear(in_features=1024, out_features=24, bias=True)
        (shared_convs): ModuleList()
        (shared_fcs): ModuleList(
            (0): Linear(in_features=12544, out_features=1024, bias=True)
            (1): Linear(in_features=1024, out_features=1024, bias=True)
        )
        (cls_convs): ModuleList()
        (cls_fcs): ModuleList()
        (reg_convs): ModuleList()
        (reg_fcs): ModuleList()
        (relu): ReLU(inplace=True)
    )
)
init_cfg=[{'type': 'Normal', 'std': 0.01, 'override': {'name': 'fc_cls'}}, {'type': 'Normal', 'std': 0.001, 'override': {'name': 'fc_reg'}}, {'type': 'Xavier', 'distribution': 'uniform', 'override': [{'name': 'shared_fcs'}, {'name': 'cls_fcs'}, {'name': 'reg_fcs'}]}]
(mask_roi_extractor): SingleROIExtractor(
    (roi_layers): ModuleList(
        (0): RoIAlign(output_size=(14, 14), spatial_scale=0.25, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
        (1): RoIAlign(output_size=(14, 14), spatial_scale=0.125, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
        (2): RoIAlign(output_size=(14, 14), spatial_scale=0.0625, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
        (3): RoIAlign(output_size=(14, 14), spatial_scale=0.03125, sampling_ratio=0, pool_mode=avg, aligned=True, use_torchvision=False)
    )
)

```

```
)  
    (mask_head): FCNMaskHead(  
        (loss_mask): CrossEntropyLoss(avg_non_ignore=False)  
        (convs): ModuleList(  
            (0): ConvModule(  
                (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
                (activate): ReLU(inplace=True)  
            )  
            (1): ConvModule(  
                (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
                (activate): ReLU(inplace=True)  
            )  
            (2): ConvModule(  
                (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
                (activate): ReLU(inplace=True)  
            )  
            (3): ConvModule(  
                (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
                (activate): ReLU(inplace=True)  
            )  
        )  
        (upsample): ConvTranspose2d(256, 256, kernel_size=(2, 2), stride=(2, 2))  
        (conv_logits): Conv2d(256, 6, kernel_size=(1, 1), stride=(1, 1))  
        (relu): ReLU(inplace=True)  
    )  
)
```

## 6.3 test dataによるLaboro Tomato 事前学習modelの正当性確認

In [12]:

```
# single-gpu testing
!python tools/test.py configs/mask_rcnn/mask_rcnn_r50_fpn_1x_coco.py \
    checkpoints/laboro_tomato_48ep.pth --show --eval "bbox" "segm"

/content/mmdetection/mmdet/utils/setup_env.py:33: UserWarning: Setting OMP_NUM_THREADS
environment variable for each process to be 1 in default, to avoid your system being o
verloaded, please further tune the variable for optimal performance in your applicatio
n as needed.
f'Setting OMP_NUM_THREADS environment variable for each process '
/content/mmdetection/mmdet/utils/setup_env.py:43: UserWarning: Setting MKL_NUM_THREADS
environment variable for each process to be 1 in default, to avoid your system being o
verloaded, please further tune the variable for optimal performance in your applicatio
n as needed.
f'Setting MKL_NUM_THREADS environment variable for each process '
loading annotations into memory...
Done (t=0.04s)
creating index...
index created!
load checkpoint from local path: checkpoints/laboro_tomato_48ep.pth
[                               ] 0/161, elapsed: 0s, ETA:/usr/loca
l/lib/python3.7/dist-packages/torch/nn/functional.py:718: UserWarning: Named tensors a
nd all their associated APIs are an experimental feature and subject to change. Please
do not use them for anything important until they are released as stable. (Triggered i
nternally at /pytorch/c10/core/TensorImpl.h:1156.)
    return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)
<Figure size 3024.01x4032.01 with 1 Axes>
[ ] 1/161, 0.1 task/s, elapsed: 9s, ETA: 1431s<Figure size 3120.01x4160.01 with 1 Ax
es>
[ ] 2/161, 0.2 task/s, elapsed: 12s, ETA: 981s<Figure size 3024.01x4032.01 with 1 A
xes>
[ ] 3/161, 0.2 task/s, elapsed: 18s, ETA: 956s<Figure size 3120.01x4160.01 with 1 A
```

```
yes>
[ ] 4/161, 0.2 task/s, elapsed: 23s, ETA: 914s<Figure size 3120.01x4160.01 with 1 A
yes>
[ ] 5/161, 0.2 task/s, elapsed: 26s, ETA: 819s<Figure size 3120.01x4160.01 with 1 A
yes>
[ ] 6/161, 0.2 task/s, elapsed: 28s, ETA: 735s<Figure size 3120.01x4160.01 with 1 A
yes>
[ ] 7/161, 0.2 task/s, elapsed: 33s, ETA: 726s<Figure size 3024.01x4032.01 with 1 A
yes>
[ ] 8/161, 0.2 task/s, elapsed: 35s, ETA: 676s<Figure size 3024.01x4032.01 with 1 A
yes>
[ ] 9/161, 0.2 task/s, elapsed: 38s, ETA: 636s<Figure size 3120.01x4160.01 with 1 A
yes>
[ ] 10/161, 0.2 task/s, elapsed: 40s, ETA: 607s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 11/161, 0.2 task/s, elapsed: 45s, ETA: 613s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 12/161, 0.3 task/s, elapsed: 47s, ETA: 587s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 13/161, 0.3 task/s, elapsed: 50s, ETA: 565s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 14/161, 0.3 task/s, elapsed: 53s, ETA: 553s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 15/161, 0.3 task/s, elapsed: 55s, ETA: 535s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 16/161, 0.3 task/s, elapsed: 57s, ETA: 518s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 17/161, 0.3 task/s, elapsed: 61s, ETA: 513s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 18/161, 0.3 task/s, elapsed: 64s, ETA: 506s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 19/161, 0.3 task/s, elapsed: 66s, ETA: 494s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 20/161, 0.3 task/s, elapsed: 69s, ETA: 487s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 21/161, 0.3 task/s, elapsed: 74s, ETA: 495s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 22/161, 0.3 task/s, elapsed: 78s, ETA: 490s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 23/161, 0.3 task/s, elapsed: 80s, ETA: 479s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 24/161, 0.3 task/s, elapsed: 82s, ETA: 468s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 25/161, 0.3 task/s, elapsed: 84s, ETA: 458s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 26/161, 0.3 task/s, elapsed: 88s, ETA: 455s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 27/161, 0.3 task/s, elapsed: 90s, ETA: 449s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 28/161, 0.3 task/s, elapsed: 94s, ETA: 445s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 29/161, 0.3 task/s, elapsed: 96s, ETA: 436s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 30/161, 0.3 task/s, elapsed: 99s, ETA: 431s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 31/161, 0.3 task/s, elapsed: 102s, ETA: 427s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 32/161, 0.3 task/s, elapsed: 104s, ETA: 420s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 33/161, 0.3 task/s, elapsed: 107s, ETA: 415s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 34/161, 0.3 task/s, elapsed: 109s, ETA: 408s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 35/161, 0.3 task/s, elapsed: 111s, ETA: 401s<Figure size 3120.01x4160.01 with 1
```

```
Axes>
[ ] 36/161, 0.3 task/s, elapsed: 113s, ETA: 394s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 37/161, 0.3 task/s, elapsed: 117s, ETA: 391s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 38/161, 0.3 task/s, elapsed: 120s, ETA: 388s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 39/161, 0.3 task/s, elapsed: 123s, ETA: 386s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 40/161, 0.3 task/s, elapsed: 127s, ETA: 385s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 41/161, 0.3 task/s, elapsed: 131s, ETA: 382s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 42/161, 0.3 task/s, elapsed: 133s, ETA: 378s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 43/161, 0.3 task/s, elapsed: 137s, ETA: 376s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 44/161, 0.3 task/s, elapsed: 140s, ETA: 372s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 45/161, 0.3 task/s, elapsed: 144s, ETA: 370s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 46/161, 0.3 task/s, elapsed: 150s, ETA: 376s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 47/161, 0.3 task/s, elapsed: 153s, ETA: 371s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 48/161, 0.3 task/s, elapsed: 156s, ETA: 366s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 49/161, 0.3 task/s, elapsed: 159s, ETA: 363s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 50/161, 0.3 task/s, elapsed: 163s, ETA: 362s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 51/161, 0.3 task/s, elapsed: 167s, ETA: 361s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 52/161, 0.3 task/s, elapsed: 170s, ETA: 357s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 53/161, 0.3 task/s, elapsed: 176s, ETA: 358s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 54/161, 0.3 task/s, elapsed: 179s, ETA: 355s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 55/161, 0.3 task/s, elapsed: 189s, ETA: 364s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 56/161, 0.3 task/s, elapsed: 192s, ETA: 359s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 57/161, 0.3 task/s, elapsed: 196s, ETA: 358s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 58/161, 0.3 task/s, elapsed: 200s, ETA: 356s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 59/161, 0.3 task/s, elapsed: 210s, ETA: 363s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 60/161, 0.3 task/s, elapsed: 218s, ETA: 366s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 61/161, 0.3 task/s, elapsed: 220s, ETA: 361s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 62/161, 0.3 task/s, elapsed: 223s, ETA: 356s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 63/161, 0.3 task/s, elapsed: 227s, ETA: 353s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 64/161, 0.3 task/s, elapsed: 230s, ETA: 348s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 65/161, 0.3 task/s, elapsed: 237s, ETA: 350s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 66/161, 0.3 task/s, elapsed: 241s, ETA: 346s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 67/161, 0.3 task/s, elapsed: 243s, ETA: 341s<Figure size 3120.01x4160.01 with 1
```

```
Axes>
[ ] 68/161, 0.3 task/s, elapsed: 247s, ETA: 338s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 69/161, 0.3 task/s, elapsed: 250s, ETA: 334s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 70/161, 0.3 task/s, elapsed: 253s, ETA: 330s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 71/161, 0.3 task/s, elapsed: 261s, ETA: 330s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 72/161, 0.3 task/s, elapsed: 268s, ETA: 331s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 73/161, 0.3 task/s, elapsed: 273s, ETA: 329s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 74/161, 0.3 task/s, elapsed: 278s, ETA: 326s<Figure size 4160.01x3120.01 with 1
Axes>
[ ] 75/161, 0.3 task/s, elapsed: 284s, ETA: 326s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 76/161, 0.3 task/s, elapsed: 292s, ETA: 326s<Figure size 4160.01x3120.01 with 1
Axes>
[ ] 77/161, 0.3 task/s, elapsed: 296s, ETA: 323s<Figure size 3120.01x4160.01 with 1
Axes>
[ ] 78/161, 0.3 task/s, elapsed: 302s, ETA: 321s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 79/161, 0.3 task/s, elapsed: 309s, ETA: 321s<Figure size 3024.01x4032.01 with 1
Axes>
[ ] 80/161, 0.3 task/s, elapsed: 311s, ETA: 315s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 81/161, 0.3 task/s, elapsed: 314s, ETA: 310s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 82/161, 0.3 task/s, elapsed: 317s, ETA: 305s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 83/161, 0.3 task/s, elapsed: 322s, ETA: 302s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 84/161, 0.3 task/s, elapsed: 329s, ETA: 302s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 85/161, 0.3 task/s, elapsed: 334s, ETA: 299s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 86/161, 0.3 task/s, elapsed: 339s, ETA: 296s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 87/161, 0.3 task/s, elapsed: 343s, ETA: 291s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 88/161, 0.3 task/s, elapsed: 349s, ETA: 289s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 89/161, 0.3 task/s, elapsed: 352s, ETA: 285s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 90/161, 0.3 task/s, elapsed: 359s, ETA: 283s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 91/161, 0.2 task/s, elapsed: 365s, ETA: 281s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 92/161, 0.2 task/s, elapsed: 369s, ETA: 277s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 93/161, 0.2 task/s, elapsed: 372s, ETA: 272s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 94/161, 0.2 task/s, elapsed: 377s, ETA: 268s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 95/161, 0.2 task/s, elapsed: 383s, ETA: 266s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 96/161, 0.2 task/s, elapsed: 387s, ETA: 262s<Figure size 3120.01x4160.01 with 1
Axes>
[> ] 97/161, 0.2 task/s, elapsed: 391s, ETA: 258s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 98/161, 0.2 task/s, elapsed: 394s, ETA: 254s<Figure size 3024.01x4032.01 with 1
Axes>
[> ] 99/161, 0.2 task/s, elapsed: 397s, ETA: 249s<Figure size 3120.01x4160.01 with 1
```

```
Axes>
[> ] 100/161, 0.2 task/s, elapsed: 401s, ETA: 244s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 101/161, 0.3 task/s, elapsed: 403s, ETA: 239s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 102/161, 0.2 task/s, elapsed: 409s, ETA: 237s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 103/161, 0.3 task/s, elapsed: 412s, ETA: 232s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 104/161, 0.2 task/s, elapsed: 418s, ETA: 229s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 105/161, 0.2 task/s, elapsed: 421s, ETA: 225s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 106/161, 0.3 task/s, elapsed: 424s, ETA: 220s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 107/161, 0.3 task/s, elapsed: 426s, ETA: 215s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 108/161, 0.3 task/s, elapsed: 430s, ETA: 211s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 109/161, 0.3 task/s, elapsed: 432s, ETA: 206s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 110/161, 0.3 task/s, elapsed: 437s, ETA: 202s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 111/161, 0.3 task/s, elapsed: 442s, ETA: 199s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 112/161, 0.3 task/s, elapsed: 445s, ETA: 195s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 113/161, 0.3 task/s, elapsed: 450s, ETA: 191stcmalloc: large alloc 1219280896 b
ytes == 0x14331a000 @ 0x7fed4f8a21e7 0x7fed4d1b30ce 0x7fed4d209cf5 0x7fed4d2b286d 0x7
fed4d2b317f 0x7fed4d2b32d0 0x4bc4ab 0x7fed4d1f4944 0x59371f 0x515244 0x549576 0x593fce
0x548ae9 0x5127f1 0x549e0e 0x4bca8a 0x7fed4d1f4944 0x59371f 0x515244 0x549576 0x593fce
0x548ae9 0x5127f1 0x549576 0x593fce 0x548ae9 0x5127f1 0x549576 0x593fce 0x5118f8 0x593
dd7
<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 114/161, 0.2 task/s, elapsed: 470s, ETA: 194s<Figure size 4160.01x3120.01 with
1 Axes>
[> ] 115/161, 0.2 task/s, elapsed: 476s, ETA: 191s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 116/161, 0.2 task/s, elapsed: 481s, ETA: 187s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 117/161, 0.2 task/s, elapsed: 486s, ETA: 183s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 118/161, 0.2 task/s, elapsed: 490s, ETA: 178s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 119/161, 0.2 task/s, elapsed: 492s, ETA: 174s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 120/161, 0.2 task/s, elapsed: 494s, ETA: 169s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 121/161, 0.2 task/s, elapsed: 497s, ETA: 164s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 122/161, 0.2 task/s, elapsed: 501s, ETA: 160s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 123/161, 0.2 task/s, elapsed: 508s, ETA: 157s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 124/161, 0.2 task/s, elapsed: 512s, ETA: 153s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 125/161, 0.2 task/s, elapsed: 517s, ETA: 149s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 126/161, 0.2 task/s, elapsed: 526s, ETA: 146s<Figure size 3024.01x4032.01 with
1 Axes>
[> ] 127/161, 0.2 task/s, elapsed: 531s, ETA: 142s<Figure size 3120.01x4160.01 with
1 Axes>
[> ] 128/161, 0.2 task/s, elapsed: 534s, ETA: 138s<Figure size 3024.01x4032.01 with
1 Axes>
```

```
[> ] 129/161, 0.2 task/s, elapsed: 537s, ETA: 133s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 130/161, 0.2 task/s, elapsed: 542s, ETA: 129s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 131/161, 0.2 task/s, elapsed: 545s, ETA: 125s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 132/161, 0.2 task/s, elapsed: 548s, ETA: 120s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 133/161, 0.2 task/s, elapsed: 551s, ETA: 116s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 134/161, 0.2 task/s, elapsed: 557s, ETA: 112s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 135/161, 0.2 task/s, elapsed: 560s, ETA: 108s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 136/161, 0.2 task/s, elapsed: 567s, ETA: 104s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 137/161, 0.2 task/s, elapsed: 576s, ETA: 101s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 138/161, 0.2 task/s, elapsed: 582s, ETA: 97s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 139/161, 0.2 task/s, elapsed: 586s, ETA: 93s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 140/161, 0.2 task/s, elapsed: 598s, ETA: 90s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 141/161, 0.2 task/s, elapsed: 600s, ETA: 85s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 142/161, 0.2 task/s, elapsed: 604s, ETA: 81s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 143/161, 0.2 task/s, elapsed: 610s, ETA: 77s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 144/161, 0.2 task/s, elapsed: 613s, ETA: 72s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 145/161, 0.2 task/s, elapsed: 617s, ETA: 68s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 146/161, 0.2 task/s, elapsed: 621s, ETA: 64s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 147/161, 0.2 task/s, elapsed: 625s, ETA: 60s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 148/161, 0.2 task/s, elapsed: 630s, ETA: 55s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 149/161, 0.2 task/s, elapsed: 636s, ETA: 51s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 150/161, 0.2 task/s, elapsed: 640s, ETA: 47s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 151/161, 0.2 task/s, elapsed: 644s, ETA: 43s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 152/161, 0.2 task/s, elapsed: 652s, ETA: 39s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 153/161, 0.2 task/s, elapsed: 657s, ETA: 34s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 154/161, 0.2 task/s, elapsed: 662s, ETA: 30s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 155/161, 0.2 task/s, elapsed: 665s, ETA: 26s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 156/161, 0.2 task/s, elapsed: 670s, ETA: 21s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 157/161, 0.2 task/s, elapsed: 675s, ETA: 17s<Figure size 3024.01x4032.01 with 1 Axes>
[> ] 158/161, 0.2 task/s, elapsed: 678s, ETA: 13s<Figure size 3120.01x4160.01 with 1 Axes>
[> ] 159/161, 0.2 task/s, elapsed: 685s, ETA: 9s<Figure size 4160.01x3120.01 with 1 Axes>
[> ] 160/161, 0.2 task/s, elapsed: 691s, ETA: 4s<Figure size 3120.01x4160.01 with 1 Axes>
```

```
[>>] 161/161, 0.2 task/s, elapsed: 695s, ETA:    0s
Evaluating bbox...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=0.82s).
Accumulating evaluation results...
DONE (t=0.09s).

Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.647
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=1000 ] = 0.822
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=1000 ] = 0.735
Average Precision (AP) @[ IoU=0.50:0.95 | area= small  | maxDets=1000 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=1000 ] = 0.146
Average Precision (AP) @[ IoU=0.50:0.95 | area= large  | maxDets=1000 ] = 0.681
Average Recall   (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.742
Average Recall   (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=300 ] = 0.742
Average Recall   (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=1000 ] = 0.742
Average Recall   (AR) @[ IoU=0.50:0.95 | area= small  | maxDets=1000 ] = 0.000
Average Recall   (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=1000 ] = 0.186
Average Recall   (AR) @[ IoU=0.50:0.95 | area= large  | maxDets=1000 ] = 0.780

Evaluating segm...
/content/mmdetection/mmdet/datasets/coco.py:474: UserWarning: The key "bbox" is deleted for more accurate mask AP of small/medium/large instances since v2.12.0. This does not change the overall mAP calculation.
  UserWarning)
Loading and preparing results...
DONE (t=0.05s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *segm*
DONE (t=1.05s).

Accumulating evaluation results...
/usr/local/lib/python3.7/dist-packages/pycocotools/cocoeval.py:378: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here.
  DeprecationWarning)
Deprecation in NumPy 1.20: for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
  tp_sum = np.cumsum(tps, axis=1).astype(dtype=np.float)
DONE (t=0.10s).

Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.660
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=1000 ] = 0.818
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=1000 ] = 0.736
Average Precision (AP) @[ IoU=0.50:0.95 | area= small  | maxDets=1000 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=1000 ] = 0.131
Average Precision (AP) @[ IoU=0.50:0.95 | area= large  | maxDets=1000 ] = 0.697
Average Recall   (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.751
Average Recall   (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=300 ] = 0.751
Average Recall   (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=1000 ] = 0.751
Average Recall   (AR) @[ IoU=0.50:0.95 | area= small  | maxDets=1000 ] = 0.000
Average Recall   (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=1000 ] = 0.187
Average Recall   (AR) @[ IoU=0.50:0.95 | area= large  | maxDets=1000 ] = 0.791

OrderedDict([('bbox_mAP', 0.647), ('bbox_mAP_50', 0.822), ('bbox_mAP_75', 0.735), ('bbox_mAP_s', 0.0), ('bbox_mAP_m', 0.146), ('bbox_mAP_l', 0.681), ('bbox_mAP_copypaste', '0.647 0.822 0.735 0.000 0.146 0.681'), ('segm_mAP', 0.66), ('segm_mAP_50', 0.818),
```

```
('segm_mAP_75', 0.736), ('segm_mAP_s', 0.0), ('segm_mAP_m', 0.131), ('segm_mAP_l', 0.697), ('segm_mAP_copypaste', '0.660 0.818 0.736 0.000 0.131 0.697')])
```

## 6.4 新たに準備したデータによる検証 (verification)

In [13]:

```
# フリー素材写真での評価（トマト、ミニトマト8枚、リンゴ：2枚）
# 前処理：写真の幅を360ピクセルに揃え、./data/eval_tomatoにロード
from sys import argv
import os
import glob
from PIL import Image

width = 640

src = glob.glob('/content/drive/MyDrive/Colab Notebooks/product_develop/eval_tomato/*')

!mkdir -p './data/eval_tomato/'
dst = './data/eval_tomato/' # リサイズ画像の保存フォルダ

for f in src:
    img = Image.open(f)
    original_width, original_height = img.size
    scale = width / original_width
    height = int(original_height * scale)
    img = img.resize((width, height))
    img.save(dst + os.path.basename(f))
```

In [14]:

```
# 事前学習済みDetectorによる推定結果の評価
import mmcv
import matplotlib.pyplot as plt
import os
import glob
import pickle

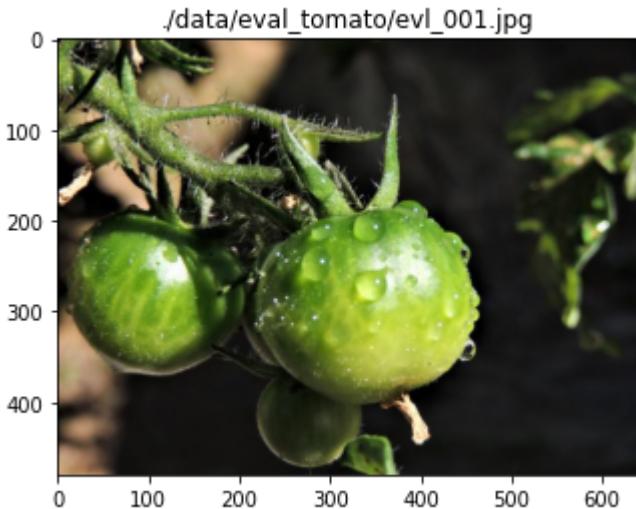
!mkdir -p './data/eval_result/'
dst = './data/eval_result/'
src_imgs = sorted(glob.glob('./data/eval_tomato/*.jpg')) # オリジナル画像のパスと拡張名

results = []

for i, f in enumerate(src_imgs):
    img = mmcv.imread(f)
    # plt.figure(figsize=(10, 10))
    plt.title(str(f))
    plt.imshow(mmcv.bgr2rgb(img))
    plt.show()

    # 検出器での物体検出推定実行
    # # Use the detector to do inference
    result = inference_detector(model, img)
    # Let's plot the result
    show_result_pyplot(model, img, result, score_thr=0.3)
    # save result data
    results.append(result)

# 推定・mask結果ファイルをpickle形式で保存
with open(dst+'results.bin', 'wb') as p:
    pickle.dump(dst+'results.bin', p)
```

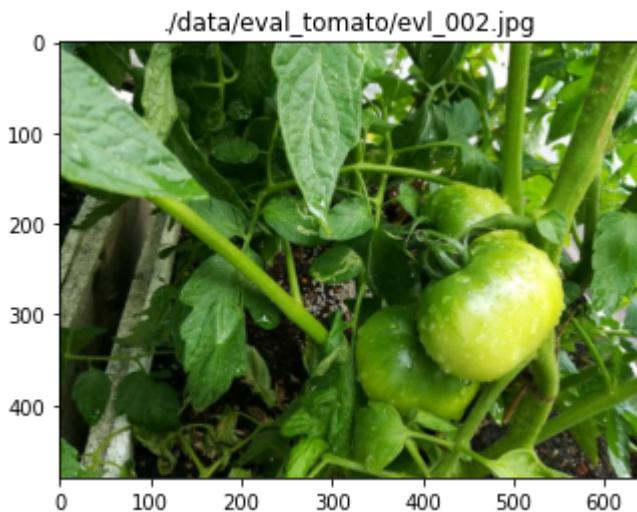


/content/mmdetection/mmdet/datasets/utils.py:70: UserWarning: "ImageToTensor" pipeline is replaced by "DefaultFormatBundle" for batch inference. It is recommended to manually replace it in the test data pipeline in your config file.

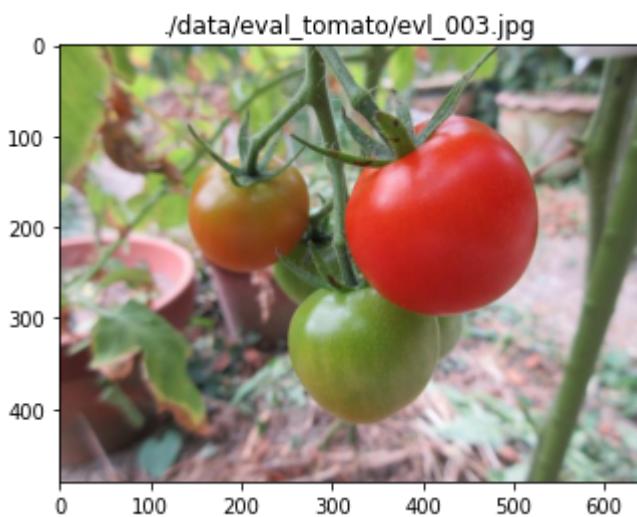
'data pipeline in your config file.', UserWarning)  
/usr/local/lib/python3.7/dist-packages/torch/nn/functional.py:718: UserWarning: Named tensors and all their associated APIs are an experimental feature and subject to change. Please do not use them for anything important until they are released as stable. (Triggered internally at /pytorch/c10/core/TensorImpl.h:1156.)  
return torch.max\_pool2d(input, kernel\_size, stride, padding, dilation, ceil\_mode)

result





result



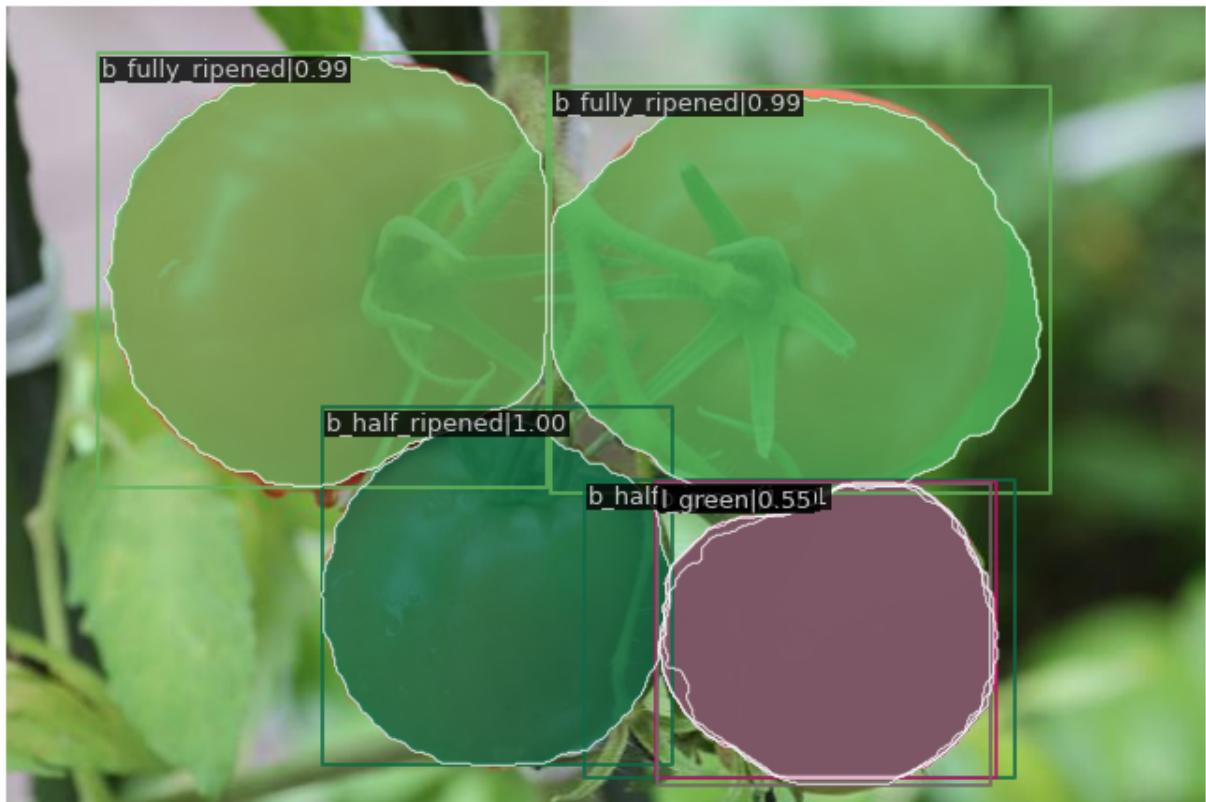
result



./data/eval\_tomato/evl\_004.jpg



result



./data/eval\_tomato/evl\_005.jpg



result



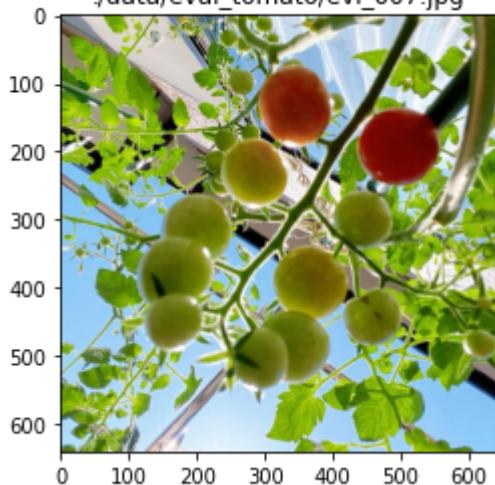
./data/eval\_tomato/evl\_006.jpg



result



./data/eval\_tomato/evl\_007.jpg



## result



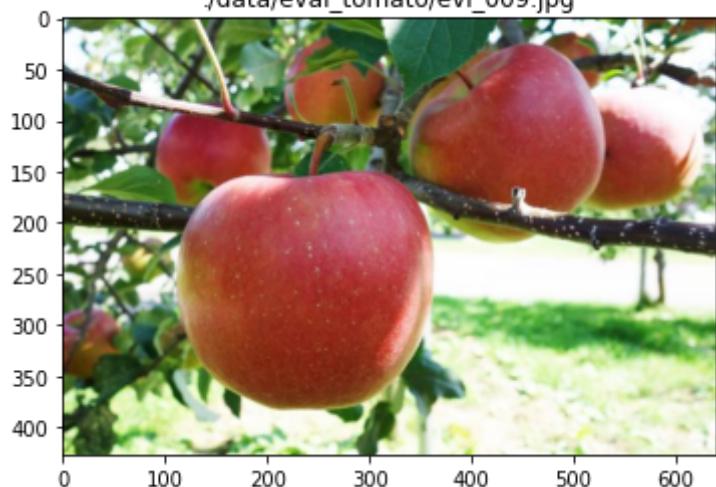
./data/eval\_tomato/evl\_008.jpg



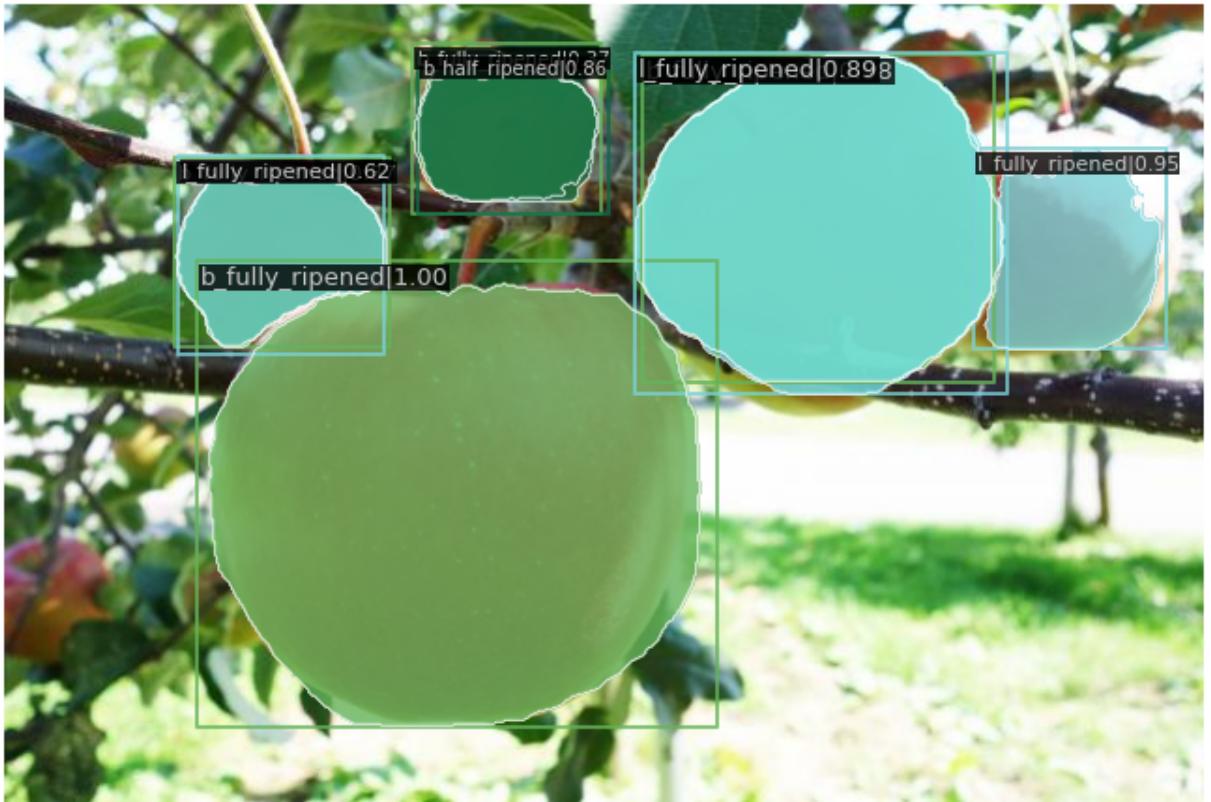
result



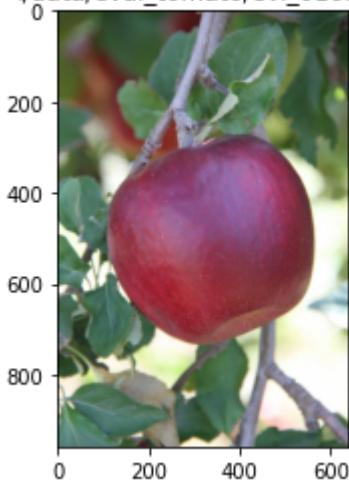
./data/eval\_tomato/evl\_009.jpg



result



./data/eval\_tomato/evl\_010.jpg



result



## 4.2 動画のセグメンテーション

In [ ]:

```
# 動画のセグメンテーション  
# 動画データのロード
```

```
! cd '/content/mmdetection'

# laboro_tomato datasetを./data/にシンボリックリンク
! mkdir -p ./data
! cp -Ri '/content/drive/MyDrive/Colab Notebooks/product_develop/video_tomato' './data'
```

In [18]:

```
# 動画セグメンテーション実行
src_mvos = sorted(glob.glob('./data/video_tomato/*.mp4')) # オリジナル画像のパスと拡張名

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

#!mkdir -p ./data/video_tomato/
#dst = './data/video_tomato' # リサイズ画像の保存フォルダ
#!ln -s '/content/drive/MyDrive/Colab Notebooks/product_develop/video_tomato' './data'

!python "./demo/video_demo.py" "./data/video_tomato/tomato3.mp4" \
    "./configs/mask_rcnn/mask_rcnn_r50_fpn_1x_coco.py" \
    "./checkpoints/laboro_tomato_48ep.pth" \
    --out "./data/eval_result/tomato3.mp4"
```

```
load checkpoint from local path: ./checkpoints/laboro_tomato_48ep.pth
[           ] 0/1752, elapsed: 0s, ETA:/content/mmdetection/mmdet/datasets/utils.py:70: UserWarning: "ImageToTensor" pipeline is replaced by "DefaultFormatBundle" for batch inference. It is recommended to manually replace it in the test data pipeline in your config file.
'data pipeline in your config file.', UserWarning)
/usr/local/lib/python3.7/dist-packages/torch/nn/functional.py:718: UserWarning: Named tensors and all their associated APIs are an experimental feature and subject to change. Please do not use them for anything important until they are released as stable. (Triggered internally at /pytorch/c10/core/TensorImpl.h:1156.)
    return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)
[>>] 1752/1752, 5.3 task/s, elapsed: 333s, ETA:      0s
```

In [ ]: